

VISVESVARAYA TECHNOLOGICAL UNIVERSITY

“JnanaSangama”, Belgaum- 590014, Karnataka.



LAB RECORD

on

Big Data Analytics (23CS6PCBDA)

Submitted by

Vedanth CR(1BM23CS036)

in partial fulfillment for the award of the degree of

BACHELOR OF ENGINEERING

in

COMPUTER SCIENCE AND ENGINEERING



B.M.S. COLLEGE OF ENGINEERING

(Autonomous Institution under VTU)

BENGALURU - 560019

January 2026 – May 2026

B.M.S. College of Engineering

Bull Temple Road, Bangalore 560019

(Affiliated to Visvesvaraya Technological University, Belgaum)

Department of Computer Science and Engineering



CERTIFICATE

This is to certify that the Lab work entitled “Big Data Analytics” carried out by **Vedanth CR (1BM23CS036)**, who is bonafide student of **B.M.S. College of Engineering**. It is in partial fulfilment for the award of **Bachelor of Engineering in Computer Science and Engineering** of the Visvesvaraya Technological University, Belgaum during the year 2026. The Lab report has been approved as it satisfies the academic requirements in respect of a Big Data Analytics (23CS6PCBDA) work prescribed for the said degree.

Anita H K

Assistant Professor Department
of CSE, BMSCE

Dr. Kavitha

Sooda Professor & HOD
Department of CSE,
BMSCE

INDEX

Sl. No.	Date	Experiment Title	Page No.
1	02.02.26	MongoDB- CRUD Operations Demonstration (Practice and Self Study)	1-8
2	09.02.26	Perform the following DB operations using Cassandra. • Create a keyspace by name Employee • Create a column family by name ○ Employee-Info with attributes ○ Emp_Id Primary Key, Emp_Name, Designation, ○ Date_of_Joining, Salary, Dept_Name • Insert the values into the table in batch • Update Employee name and Department of EmpId 121 • Sort the details of Employee records based on salary • Alter the schema of the table Employee_Info to add a column Projects which stores a set of Projects done by the corresponding Employee. • Update the altered table to add project names. • Create a TTL of 15 seconds to display the values of Employees.	9-10
3	16.02.26	Perform the following DB operations using Cassandra. • Create a keyspace by name Library • Create a column family by name Library-Info with attributes ○ Stud_Id Primary Key, ○ Counter_value of type Counter, ○ Stud_Name, Book-Name, Book-Id, ○ Date_of_issue • Insert the values into the table in batch • Display the details of the table created and increase the value of the counter • Write a query to show that a student with id 112 has taken a book “BDA” 2 times. • Export the created column to a csv file • g) Import a given csv dataset from local file system into Cassandra column family	11-14
4	09.03.26	Execution of HDFS Commands for interaction with Hadoop Environment. (Minimum 10 commands to be executed)	15-16
5	09.03.26	Implement Wordcount program on Hadoop framework	17-21
6	16.03.26	From the following link extract the weather data https://github.com/tomwhite/hadoop-book/tree/master/input/ncdc/all • Create a MapReduce program to find average temperature for each year from NCDC data set. • b) find the mean max temperature for every month	22-27

7	23.03.26	For a given Text file, Create a Map Reduce program to sort the content in an alphabetic order listing only top 10 maximum occurrences of words.	28-32
8	09.04.26	Write a Scala program to print numbers from 1 to 100 using for loop.	33
9	09.04.26	Using RDD and FlatMap count how many times each word appears in a file and write out a list of words whose count is strictly greater than 4 using Spark.	34
10	16.04.26	Write a simple streaming program in Spark to receive text data streams on a particular port, perform basic text cleaning (like white space removal, stop words removal, lemmatization, etc.), and print the cleaned text on the screen. (Open Ended Question).	35

Github Link: <https://github.com/vedanthcr27/BDALAB>

Course Outcomes (COs):

CO1	Apply the concept of NoSQL, Hadoop or Spark for a given task
CO2	Analyse big data analytics mechanisms that can be applied to obtain solution for a given problem.
CO3	Design and implement solutions using data analytics mechanisms for a given problem.

PROGRAM – 1

MongoDB- CRUD Operations Demonstration (Practice and Self Study)

Commands with output:

Student database

1. Create a database “Student” with the following attributes Rollno, Age, ContactNo,

```
Atlas atlas-b6pfyk-shard-0 [primary] test> db.createCollection("Student");
{ ok: 1 }
```

Email-Id.

2. Insert appropriate values (at least 5)

```
Atlas atlas-b6pfyk-shard-0 [primary] test> db.Student.insert({RollNo:1, Age:21, Cont:9876, email:"antara.de9@gmail.com"});
DeprecationWarning: Collection.insert() is deprecated. Use insertOne, insertMany, or bulkWrite.
{
  acknowledged: true,
  insertedIds: { '0': ObjectId("660a82ec7c840f42b4a05530") }
}
Atlas atlas-b6pfyk-shard-0 [primary] test>

Atlas atlas-b6pfyk-shard-0 [primary] test> db.Student.insert({RollNo:2, Age:22, Cont:9976, email:"anushka.de9@gmail.com"});
{
  acknowledged: true,
  insertedIds: { '0': ObjectId("660a82ed7c840f42b4a05531") }
}
Atlas atlas-b6pfyk-shard-0 [primary] test>

Atlas atlas-b6pfyk-shard-0 [primary] test> db.Student.insert({RollNo:3, Age:21, Cont:5576, email:"anubhav.de9@gmail.com"});
{
  acknowledged: true,
  insertedIds: { '0': ObjectId("660a82ed7c840f42b4a05532") }
}
Atlas atlas-b6pfyk-shard-0 [primary] test>

Atlas atlas-b6pfyk-shard-0 [primary] test> db.Student.insert({RollNo:4, Age:20, Cont:4476, email:"pani.de9@gmail.com"});
{
  acknowledged: true,
  insertedIds: { '0': ObjectId("660a82ed7c840f42b4a05533") }
}
Atlas atlas-b6pfyk-shard-0 [primary] test>

Atlas atlas-b6pfyk-shard-0 [primary] test> db.Student.insert({RollNo:10, Age:23, Cont:2276, email:"rekha.de9@gmail.com"});
{
  acknowledged: true,
```

```
Atlas atlas-b6pfyk-shard-0 [primary] test> db.Student.insert({RollNo:10, Age:23, Cont:2276, email:"rekha.de9@gmail.com"});
{
  acknowledged: true,
  insertedIds: { '0': ObjectId("660a82f47c840f42b4a05534") }
}
```

3. View the data

```
Atlas atlas-b6pfyk-shard-0 [primary] test> db.Student.find()
[
  {
    _id: ObjectId("660a82ec7c840f42b4a05530"),
    RollNo: 1,
    Age: 21,
    Cont: 9876,
    email: 'antara.de9@gmail.com'
  },
  {
    _id: ObjectId("660a82ed7c840f42b4a05531"),
    RollNo: 2,
    Age: 22,
    Cont: 9976,
    email: 'anushka.de9@gmail.com'
  },
  {
    _id: ObjectId("660a82ed7c840f42b4a05532"),
    RollNo: 3,
    Age: 21,
    Cont: 5576,
    email: 'anubhav.de9@gmail.com'
  },
  {
    _id: ObjectId("660a82ed7c840f42b4a05533"),
    RollNo: 4,
    Age: 20,
    Cont: 4476,
    email: 'pani.de9@gmail.com'
  },
  {
    _id: ObjectId("660a82f47c840f42b4a05534"),
    RollNo: 10,
    Age: 23,
    Cont: 2276,
    email: 'rekha.de9@gmail.com'
  }
]
```

4. Write a query to update the Email-Id of a student with rollno 10.

```
{
  _id: ObjectId("660a83337c840f42b4a05535"),
  RollNo: 11,
  Age: 22,
  Name: 'ABC',
  Cont: 2276,
  email: 'rea.de9@gmail.com'
}
```

5. Replace the student name from “ABC” to “FEM” of rollno 11.

```
Atlas atlas-b6pfyk-shard-0 [primary] test> db.Student.update({RollNo:11,Name:"ABC"},{$set:{Name:"FEM"}})
{
  acknowledged: true,
  insertedId: null,
  matchedCount: 1,
  modifiedCount: 1,
  upsertedCount: 0
}
```

6. Export the created table into local file system

```
C:\Users\BMSCECSE>mongoexport mongodb+srv://antararc:Test1234@cluster0.mfnfeys.mongodb.net/myDB --collection=Student --out C:\Users\BMSCECSE\Downloads\output.json
2023-01-12T15:15:56.383+0530 connected to: mongodb+srv://[**REDACTED**]@cluster0.mfnfeys.mongodb.net/myDB
2023-01-12T15:15:56.497+0530 exported 7 records
```

7. Drop the table

```
Atlas atlas-b6pfyk-shard-0 [primary] test> db.Student.drop();
true
```

8. Import a given csv dataset from local file system into mongodb collection.

```
C:\Users\BMSCECSE>mongoimport mongodb+srv://antararc:Test1234@cluster0.mfnfeys.mongodb.net/myDB --collection=New_Student --type json --file C:\Users\BMSCECSE\Downloads\output.json
2023-01-12T15:17:35.523+0530 connected to: mongodb+srv://[**REDACTED**]@cluster0.mfnfeys.mongodb.net/myDB
2023-01-12T15:17:35.640+0530 7 document(s) imported successfully. 0 document(s) failed to import.
```

MongoDB- CRUD Operations Demonstration (Practice and Self Study)

Command with output using terminal:

```
test> use Employee
switched to db Employee
Employee> db
Employee
Employee> db.createCollection("Employee_Info")
{ ok: 1 }
Employee> show collections
Employee_Info
Employee> db.Employee_Info.insertMany([
...   {
...     Emp_Id: 121,
...     Emp_Name: "Amit",
...     Designation: "Software Engineer",
...     Date_of_Joining: ISODate("2019-06-15"),
...     Salary: 60000,
...     Dept_Name: "IT"
...   },
...   {
...     Emp_Id: 122,
...     Emp_Name: "Ravi",
...     Designation: "Senior Engineer",
...     Date_of_Joining: ISODate("2018-03-20"),
...     Salary: 75000,
...     Dept_Name: "R&D"
...   },
...   {
...     Emp_Id: 123,
...     Emp_Name: "Meera",
...     Designation: "Manager",
...     Date_of_Joining: ISODate("2016-01-10"),
...     Salary: 90000,
...     Dept_Name: "HR"
...   }
... ])
...
{
  acknowledged: true,
  insertedIds: {
    '0': ObjectId('69899e77e7bd2cd77163b112'),
    '1': ObjectId('69899e77e7bd2cd77163b113'),
    '2': ObjectId('69899e77e7bd2cd77163b114')
  }
}
```



```
}
Employee> db.Employee_Info.find({ Emp_Id: 121 })
[
  {
    _id: ObjectId('69899e77e7bd2cd77163b112'),
    Emp_Id: 121,
    Emp_Name: 'Amit Kumar',
    Designation: 'Software Engineer',
    Date_of_Joining: ISODate('2019-06-15T00:00:00.000Z'),
    Salary: 60000,
    Dept_Name: 'Data Science'
  }
]
Employee> db.Employee_Info.find().sort({ Salary: -1 })
[
  {
    _id: ObjectId('69899e77e7bd2cd77163b114'),
    Emp_Id: 123,
    Emp_Name: 'Meera',
    Designation: 'Manager',
    Date_of_Joining: ISODate('2016-01-10T00:00:00.000Z'),
    Salary: 90000,
    Dept_Name: 'HR'
  },
  {
    _id: ObjectId('69899e77e7bd2cd77163b113'),
    Emp_Id: 122,
    Emp_Name: 'Ravi',
    Designation: 'Senior Engineer',
    Date_of_Joining: ISODate('2018-03-20T00:00:00.000Z'),
    Salary: 75000,
    Dept_Name: 'R&D'
  },
  {
    _id: ObjectId('69899e77e7bd2cd77163b112'),
    Emp_Id: 121,
    Emp_Name: 'Amit Kumar',
    Designation: 'Software Engineer',
    Date_of_Joining: ISODate('2019-06-15T00:00:00.000Z'),
    Salary: 60000,
    Dept_Name: 'Data Science'
  }
]
Employee> db.Employee_Info.find().sort({ Salary: 1 })
[
  {
    _id: ObjectId('69899e77e7bd2cd77163b112'),
    Emp_Id: 121,
    Emp_Name: 'Amit Kumar',
    Designation: 'Software Engineer',
    Date_of_Joining: ISODate('2019-06-15T00:00:00.000Z'),
```

mongosh mongodb://127.0.0.1

```
Designation: 'Software Engineer',  
Date_of_Joining: ISODate('2019-06-15T00:00:00.000Z'),  
Salary: 60000,  
Dept_Name: 'Data Science'
```

```
}
```

```
]
```

```
Employee> db.Employee_Info.find().sort({ Salary: -1 })
```

```
[
```

```
{
```

```
  _id: ObjectId('69899e77e7bd2cd77163b114'),  
  Emp_Id: 123,  
  Emp_Name: 'Meera',  
  Designation: 'Manager',  
  Date_of_Joining: ISODate('2016-01-10T00:00:00.000Z'),  
  Salary: 90000,  
  Dept_Name: 'HR'
```

```
},
```

```
{
```

```
  _id: ObjectId('69899e77e7bd2cd77163b113'),  
  Emp_Id: 122,  
  Emp_Name: 'Ravi',  
  Designation: 'Senior Engineer',  
  Date_of_Joining: ISODate('2018-03-20T00:00:00.000Z'),  
  Salary: 75000,  
  Dept_Name: 'R&D'
```

```
},
```

```
{
```

```
  _id: ObjectId('69899e77e7bd2cd77163b112'),  
  Emp_Id: 121,  
  Emp_Name: 'Amit Kumar',  
  Designation: 'Software Engineer',  
  Date_of_Joining: ISODate('2019-06-15T00:00:00.000Z'),  
  Salary: 60000,  
  Dept_Name: 'Data Science'
```

```
}
```

```
]
```

```
Employee> db.Employee_Info.find().sort({ Salary: 1 })
```

```
[
```

```
{
```

```
  _id: ObjectId('69899e77e7bd2cd77163b112'),  
  Emp_Id: 121,  
  Emp_Name: 'Amit Kumar',  
  Designation: 'Software Engineer',  
  Date_of_Joining: ISODate('2019-06-15T00:00:00.000Z'),  
  Salary: 60000,  
  Dept_Name: 'Data Science'
```

```
},
```

```
{
```

```
  _id: ObjectId('69899e77e7bd2cd77163b113'),  
  Emp_Id: 122,  
  Emp_Name: 'Ravi',
```

mongosh mongodb://127.0.0.1:27027/

```
Emp_Name: 'Ravi',
Designation: 'Senior Engineer',
Date_of_Joining: ISODate('2018-03-20T00:00:00.000Z'),
Salary: 75000,
Dept_Name: 'R&D'
```

```
},
{
```

```
  _id: ObjectId('69899e77e7bd2cd77163b114'),
  Emp_Id: 123,
  Emp_Name: 'Meera',
  Designation: 'Manager',
  Date_of_Joining: ISODate('2016-01-10T00:00:00.000Z'),
  Salary: 90000,
  Dept_Name: 'HR'
```

```
}
```

```
]
```

```
Employee> db.Employee_Info.updateOne(
```

```
...   { Emp_Id: 121 },
```

```
...   { $set: { Projects: ["AI-Analytics", "Cloud-Migration", "BigData-Pipeline"] } }  
... )
```

```
...
```

```
{
```

```
  acknowledged: true,  
  insertedId: null,  
  matchedCount: 1,  
  modifiedCount: 1,  
  upsertedCount: 0
```

```
}
```

```
Employee> db.Employee_Info.updateOne(
```

```
...   { Emp_Id: 121 },
```

```
...   { $addToSet: { Projects: "ML-Deployment" } }  
... )
```

```
...
```

```
{
```

```
  acknowledged: true,  
  insertedId: null,  
  matchedCount: 1,  
  modifiedCount: 1,  
  upsertedCount: 0
```

```
}
```

```
Employee> db.Employee_Info.find({ Emp_Id: 121 })
```

```
[
```

```
{
```

```
  _id: ObjectId('69899e77e7bd2cd77163b112'),  
  Emp_Id: 121,  
  Emp_Name: 'Amit Kumar',  
  Designation: 'Software Engineer',  
  Date_of_Joining: ISODate('2019-06-15T00:00:00.000Z'),  
  Salary: 60000,  
  Dept_Name: 'Data Science',  
  Projects: [
```

```

}
Employee> db.Employee_Info.find({ Emp_Id: 121 })
[
  {
    _id: ObjectId('69899e77e7bd2cd77163b112'),
    Emp_Id: 121,
    Emp_Name: 'Amit Kumar',
    Designation: 'Software Engineer',
    Date_of_Joining: ISODate('2019-06-15T00:00:00.000Z'),
    Salary: 60000,
    Dept_Name: 'Data Science',
    Projects: [
      'AI-Analytics',
      'Cloud-Migration',
      'BigData-Pipeline',
      'ML-Deployment'
    ]
  }
]
Employee> db.Employee_Info.insertOne({
...   Emp_Id: 130,
...   Emp_Name: "TemporaryUser",
...   Designation: "Intern",
...   Date_of_Joining: ISODate("2024-01-01"),
...   Salary: 15000,
...   Dept_Name: "Training",
...   createdAt: new Date()
... })
...
{
  acknowledged: true,
  insertedId: ObjectId('69899ef7e7bd2cd77163b115')
}
Employee> db.Employee_Info.createIndex(
...   { createdAt: 1 },
...   { expireAfterSeconds: 15 }
... )
...
createdAt_1
Employee> db.Employee_Info.find({ Emp_Id: 130 })
[
  {
    _id: ObjectId('69899ef7e7bd2cd77163b115'),
    Emp_Id: 130,
    Emp_Name: 'TemporaryUser',
    Designation: 'Intern',
    Date_of_Joining: ISODate('2024-01-01T00:00:00.000Z'),
    Salary: 15000,
    Dept_Name: 'Training',
    createdAt: ISODate('2026-02-09T08:46:47.854Z')
  }
]

```

PROGRAM – 2

Perform the following DB operations using Cassandra Questions:

- a) Create a keyspace by name Employee
- b) Create a column family by name
 - Employee-Info with attributes
 - Emp_Id Primary Key, Emp_Name,
 - Designation, Date_of_Joining,
 - Salary, Dept_Name
- c) Insert the values into the table in batch
- d) Update Employee name and Department of Emp-Id 121
- e) Sort the details of Employee records based on salary
- f) Alter the schema of the table Employee_Info to add a column Projects which stores a set of Projects done by the corresponding Employee.
- g) Update the altered table to add project names.
- h) Create a TTL of 15 seconds to display the values of Employees.

Command with output

Employee Database

1. Create a keyspace by name Employee

```
cqlsh:library> CREATE KEYSPACE Employee WITH REPLICATION = { 'class' : 'SimpleStrategy', 'replication_factor' : 1 };
cqlsh:library>
```

2. Create a column family by name Employee-Info with attributes Emp_Id Primary Key,

```
cqlsh:employee>
cqlsh:employee> CREATE TABLE Employee_Info (
...     Emp_Id int PRIMARY KEY,
...     Emp_Name text,
...     Designation text,
...     Date_of_Joining date,
...     Salary decimal,
...     Dept_Name text
... );
```

Emp_Name, Designation, Date_of_Joining, Salary, Dept_Name

3. Insert the values into the table in batch

```
cqlsh:employee> BEGIN BATCH
... INSERT INTO Employee_Info (Emp_Id, Emp_Name, Designation, Date_of_Joining, Salary, Dept_Name)
... VALUES (101, 'John Doe', 'Manager', '2023-01-01', 50000, 'HR');
... INSERT INTO Employee_Info (Emp_Id, Emp_Name, Designation, Date_of_Joining, Salary, Dept_Name)
... VALUES (121, 'Jane Smith', 'Developer', '2023-02-01', 60000, 'IT');
... APPLY BATCH;
cqlsh:employee> UPDATE Employee_Info SET Emp_Name = 'Jane Johnson', Dept_Name = 'Engineering' WHERE Emp_Id = 121;
cqlsh:employee> SELECT * FROM Employee_Info;
```

emp_id	date_of_joining	dept_name	designation	emp_name	salary
121	2023-02-01	Engineering	Developer	Jane Johnson	60000
101	2023-01-01	HR	Manager	John Doe	50000

(2 rows)

4. Update Employee name and Department of Emp-Id 121
5. Sort the details of Employee records based on salary


```

cqlsh:employee> paging off
Disabled Query paging.
cqlsh:employee> SELECT * FROM Employee_Info WHERE Emp_Id IN (121,101) ORDER BY Salary ALLOW FILTERING;

 emp_id | salary | date_of_joining | dept_name | designation | emp_name | projects
-----+-----+-----+-----+-----+-----+-----
    121 | 60000 | 2023-02-01 | IT | Developer | Jane Smith | {'ProjectC'}
    101 | 50000 | 2023-01-01 | HR | Manager | John Doe | {'ProjectA', 'ProjectB'}
(2 rows)

```

6. Alter the schema of the table Employee_Info to add a column Projects which stores a set of Projects done by the corresponding Employee.

7. Update the altered table to add project names.

8. Create a TTL of 15 seconds to display the values of Employees.

```

cqlsh:employee> UPDATE Employee_Info SET Projects = {'ProjectA', 'ProjectB'} WHERE Emp_Id = 101 and salary=50000;
cqlsh:employee> UPDATE Employee_Info SET Projects = {'ProjectC'} WHERE Emp_Id = 121 and salary=60000;
cqlsh:employee> select * from Employee_Info;

 emp_id | salary | date_of_joining | dept_name | designation | emp_name | projects
-----+-----+-----+-----+-----+-----+-----
    121 | 60000 | 2023-02-01 | IT | Developer | Jane Smith | {'ProjectC'}
    101 | 50000 | 2023-01-01 | HR | Manager | John Doe | {'ProjectA', 'ProjectB'}
(2 rows)

```

```

cqlsh:employee> INSERT INTO Employee_Info (Emp_Id, Emp_Name, Designation, Date_of_Joining, Salary, Dept_Name) VALUES (102, 'Jane Smith', 'Developer', '2022-06-03', 60000, 'IT') USING TTL 15;
cqlsh:employee> select ttl(Emp_Name) from Employee_Info where Emp_id=102;

 ttl(emp_name)
-----
          14
(1 rows)

```

PROGRAM – 3

Perform the following DB operations using Cassandra Questions:

- a) Create a keyspace by name Library
- b) Create a column family by name Library-Info with attributes
 - Stud_Id Primary Key,
 - Counter_value of type Counter,
 - Stud_Name, Book-Name, Book-Id,
 - Date_of_issue
- c) Insert the values into the table in batch
- d) Display the details of the table created and increase the value of the counter
- e) Write a query to show that a student with id 112 has taken a book “BDA” 2 times.
- f) Export the created column to a csv file
- g) Import a given csv dataset from local file system into Cassandra

column family Command with output

Use HELP for help.

```
cqlsh> DESCRIBE KEYSPACES;
```

```
system      system_distributed  system_traces  system_virtual_schema
system_auth  system_schema      system_views
```

```
cqlsh> CREATE KEYSPACE Library
```

```
... WITH replication = {'class': 'SimpleStrategy', 'replication_factor': 1};
```

```
cqlsh> USE Library;
```

```
cqlsh:library> CREATE TABLE Library_Info (
```

```
...     Stud_Id int,
...     Stud_Name text,
...     Book_Name text,
...     Book_Id int,
...     Date_of_issue date,
...     PRIMARY KEY (Stud_Id, Book_Id)
... );
```

```
cqlsh:library> CREATE TABLE Book_Count (
```

```
...     Stud_Id int,
...     Book_Name text,
...     Counter_value counter,
...     PRIMARY KEY (Stud_Id, Book_Name)
... );
```

```
cqlsh:library> BEGIN BATCH
```

```
... INSERT INTO Library_Info (Stud_Id, Stud_Name, Book_Name, Book_Id, Date_of_issue)
... VALUES (112, 'Rahul', 'BDA', 1, '2026-02-24');
```

```
...
```

```
... INSERT INTO Library_Info (Stud_Id, Stud_Name, Book_Name, Book_Id, Date_of_issue)
... VALUES (113, 'Anita', 'ML', 2, '2026-02-24');
```

```
... APPLY BATCH;
```

```
cqlsh:library> UPDATE Book_Count
```

```
... SET Counter_value = Counter_value + 1
... WHERE Stud_Id = 112 AND Book_Name = 'BDA';
```



```

... );
cqlsh:library> CREATE TABLE Book_Count (
...     Stud_Id int,
...     Book_Name text,
...     Counter_value counter,
...     PRIMARY KEY (Stud_Id, Book_Name)
... );
cqlsh:library> BEGIN BATCH
... INSERT INTO Library_Info (Stud_Id, Stud_Name, Book_Name, Book_Id, Date_of_issue)
... VALUES (112, 'Rahul', 'BDA', 1, '2026-02-24');
...
... INSERT INTO Library_Info (Stud_Id, Stud_Name, Book_Name, Book_Id, Date_of_issue)
... VALUES (113, 'Anita', 'ML', 2, '2026-02-24');
... APPLY BATCH;
cqlsh:library> UPDATE Book_Count
... SET Counter_value = Counter_value + 1
... WHERE Stud_Id = 112 AND Book_Name = 'BDA';
cqlsh:library>
cqlsh:library> UPDATE Book_Count
... SET Counter_value = Counter_value + 1
... WHERE Stud_Id = 112 AND Book_Name = 'BDA';
cqlsh:library> SELECT * FROM Library_Info;

```

stud_id	book_id	book_name	date_of_issue	stud_name
113	2	ML	2026-02-24	Anita
112	1	BDA	2026-02-24	Rahul

(2 rows)

```
cqlsh:library> SELECT * FROM Book_Count;
```

stud_id	book_name	counter_value
112	BDA	2

```

112 |      BDA |      2
(1 rows)
cqlsh:library> UPDATE Book_Count
... SET Counter_value = Counter_value + 1
... WHERE Stud_Id = 112 AND Book_Name = 'BDA';
cqlsh:library> SELECT Counter_value
... FROM Book_Count
... WHERE Stud_Id = 112 AND Book_Name = 'BDA';

```

counter_value
3

(1 rows)

```
cqlsh:library> exit
```

```

C:\Users\BPC\Downloads\apache-cassandra-4.1.10-bin\apache-cassandra-4.1.10\bin>docker exec -it my-cassandra cqlsh -e "COPY Library.Library_Info
TO '/tmp/library.csv' WITH HEADER = true;"
Using 7 child processes

```

```

Starting copy of library.library_info with columns [stud_id, book_id, book_name, date_of_issue, stud_name].
Processed: 2 rows; Rate:      7 rows/s; Avg. rate:      3 rows/s
2 rows exported to 1 files in 0.633 seconds.

```

```

C:\Users\BPC\Downloads\apache-cassandra-4.1.10-bin\apache-cassandra-4.1.10\bin>docker cp my-cassandra:/tmp/library.csv library.csv
Successfully copied 2.05kB to C:\Users\BPC\Downloads\apache-cassandra-4.1.10-bin\apache-cassandra-4.1.10\bin\library.csv

```

```

C:\Users\BPC\Downloads\apache-cassandra-4.1.10-bin\apache-cassandra-4.1.10\bin>docker cp library.csv my-cassandra:/tmp/library.csv
Successfully copied 2.05kB to my-cassandra:/tmp/library.csv

```

```

C:\Users\BPC\Downloads\apache-cassandra-4.1.10-bin\apache-cassandra-4.1.10\bin>docker exec -it my-cassandra cqlsh -e "COPY Library.Library_Info
FROM '/tmp/library.csv' WITH HEADER = true;"
Using 7 child processes

```

```

Starting copy of library.library_info with columns [stud_id, book_id, book_name, date_of_issue, stud_name].
Processed: 2 rows; Rate:      4 rows/s; Avg. rate:      5 rows/s
2 rows imported from 1 files in 0.379 seconds (0 skipped).

```


PROGRAM – 4

Execution of HDFS Commands for interaction with Hadoop Environment. (Minimum 10 commands to be executed)

Command with output

1. mkdir
2. ls

```
hadoop@bmscscse-HP-Elite-Tower-800-G9-Desktop-PC:~$ start-all.sh
WARNING: Attempting to start all Apache Hadoop daemons as hadoop in 10 seconds.
WARNING: This is not a recommended production deployment configuration.
WARNING: Use CTRL-C to abort.
Starting namenodes on [localhost]
Starting datanodes
Starting secondary namenodes [bmscscse-HP-Elite-Tower-800-G9-Desktop-PC]
Starting resourcemanager
Starting nodemanagers
hadoop@bmscscse-HP-Elite-Tower-800-G9-Desktop-PC:~$ hdfs dfs -mkdir /bda_hadoop
hadoop@bmscscse-HP-Elite-Tower-800-G9-Desktop-PC:~$ hadoop fs -ls /
Found 1 items
drwxr-xr-x  - hadoop supergroup          0 2024-05-13 14:37 /bda_hadoop
```

3. put

```
hadoop@bmscscse-HP-Elite-Tower-800-G9-Desktop-PC:~$ hdfs dfs -put /home/hadoop/Desktop/bda_local.txt /bda_hadoop/file.txt
hadoop@bmscscse-HP-Elite-Tower-800-G9-Desktop-PC:~$ hadoop fs -ls /bda_hadoop
Found 1 items
-rw-r--r--  1 hadoop supergroup          9 2024-05-13 14:42 /bda_hadoop/file.txt
hadoop@bmscscse-HP-Elite-Tower-800-G9-Desktop-PC:~$ hdfs dfs -cat /bda_hadoop/file.txt
Hello!!!
```

4. copyFromLocal

```
hadoop@bmscscse-HP-Elite-Tower-800-G9-Desktop-PC:~$ hdfs dfs -copyFromLocal /home/hadoop/Desktop/bda_local.txt /bda_hadoop/file_cp_local.txt
hadoop@bmscscse-HP-Elite-Tower-800-G9-Desktop-PC:~$ hdfs dfs -cat /bda_hadoop/file_cp_local.txt
Hello!!!
hadoop@bmscscse-HP-Elite-Tower-800-G9-Desktop-PC:~$
```

5. get

```
hadoop@bmscscse-HP-Elite-Tower-800-G9-Desktop-PC:~$ hdfs dfs -get /bda_hadoop/file.txt /home/hadoop/Desktop/downloaded_file.txt
hadoop@bmscscse-HP-Elite-Tower-800-G9-Desktop-PC:~$ hdfs dfs -getmerge /bda_hadoop/file.txt /bda_hadoop/file_cp_local.txt /home/hadoop/Desktop/downloaded_file.txt
hadoop@bmscscse-HP-Elite-Tower-800-G9-Desktop-PC:~$ hadoop fs -getfacl /bda_hadoop/
hadoop@bmscscse-HP-Elite-Tower-800-G9-Desktop-PC:~$ hadoop fs -getfacl /bda_hadoop/
# file: /bda_hadoop
# owner: hadoop
# group: supergroup
user::rwx
group::r-x
other::r-x
```

6. copyToLocal

```

hadoop@bmscsecse-HP-Elite-Tower-800-G9-Desktop-PC:~$ hdfs dfs -copyToLocal /bda_hadoop/file.txt /home/hadoop/Desktop
hadoop@bmscsecse-HP-Elite-Tower-800-G9-Desktop-PC:~$ hadoop fs -mv /bda_hadoop /abc
hadoop@bmscsecse-HP-Elite-Tower-800-G9-Desktop-PC:~$ hadoop fs -ls /abc
Found 2 items
-rw-r--r-- 1 hadoop supergroup          9 2024-05-13 14:42 /abc/file.txt
-rw-r--r-- 1 hadoop supergroup          9 2024-05-13 14:52 /abc/file_cp_local.txt
hadoop@bmscsecse-HP-Elite-Tower-800-G9-Desktop-PC:~$ hadoop fs -cp /hello/ /hadoop_lab
cp: '/hello/': No such file or directory
hadoop@bmscsecse-HP-Elite-Tower-800-G9-Desktop-PC:~$ 

```

7. cat

```

hadoop@bmscsecse-HP-Elite-Tower-800-G9-Desktop-PC:~$ hdfs dfs -cat /bda_hadoop/file_cp_local.txt
Hello!!!
hadoop@bmscsecse-HP-Elite-Tower-800-G9-Desktop-PC:~$ 

```

8.mv

```

hadoop@bmscsecse-HP-Elite-Tower-800-G9-Desktop-PC:~$ hadoop fs -mv /bda_hadoop /abc
hadoop@bmscsecse-HP-Elite-Tower-800-G9-Desktop-PC:~$ hadoop fs -ls /abc
Found 2 items
-rw-r--r-- 1 hadoop supergroup          9 2024-05-13 14:42 /abc/file.txt
-rw-r--r-- 1 hadoop supergroup          9 2024-05-13 14:52 /abc/file_cp_local.txt

```

9.cp

```

-rw-r--r-- 1 hadoop supergroup          9 2024-05-13 14:52 /abc/file_cp_local.txt
hadoop@bmscsecse-HP-Elite-Tower-800-G9-Desktop-PC:~$ hadoop fs -cp /hello/ /hadoop_lab
cp: '/hello/': No such file or directory
hadoop@bmscsecse-HP-Elite-Tower-800-G9-Desktop-PC:~$ 

```

PROGRAM – 5

Implement Wordcount program on Hadoop framework

Driver Code

```
import java.io.IOException;
import
org.apache.hadoop.conf.Configured;
import org.apache.hadoop.fs.Path;
import
org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapred.FileInputFormat;
import
org.apache.hadoop.mapred.FileOutputFormat;
import org.apache.hadoop.mapred.JobClient;
import org.apache.hadoop.mapred.JobConf;
import org.apache.hadoop.util.Tool;
import
org.apache.hadoop.util.ToolRunner;
public class WCDriver extends Configured implements Tool {

    public int run(String[] args) throws IOException {

        if (args.length < 2) {
            System.out.println("Please give valid
            inputs"); return -1;
        }

        JobConf conf = new JobConf(WCDriver.class);
        conf.setJobName("WordCount");

        FileInputFormat.setInputPaths(conf, new Path(args[0]));
        FileOutputFormat.setOutputPath(conf, new
        Path(args[1]));

        conf.setMapperClass(WCMapper.class);
        conf.setReducerClass(WCReducer.class
        );

        conf.setMapOutputKeyClass(Text.class);
        conf.setMapOutputValueClass(IntWritable.clas
        s);

        conf.setOutputKeyClass(Text.class);
        conf.setOutputValueClass(IntWritable.clas
        s);

        JobClient.runJob(conf)
        ; return 0;
    }

    // Main Method
    public static void main(String[] args) throws Exception {
        int exitCode = ToolRunner.run(new WCDriver(),
        args); System.out.println("Job Exit Code: " +
        exitCode);
    }
}
```

Mapper Code

```
import java.io.IOException;
import org.apache.hadoop.io.IntWritable;
import
org.apache.hadoop.io.LongWritable;
import org.apache.hadoop.io.Text;
import
```

```
org.apache.hadoop.mapred.MapReduceBase;  
import org.apache.hadoop.mapred.Mapper;  
import  
org.apache.hadoop.mapred.OutputCollector;  
import org.apache.hadoop.mapred.Reporter;
```

```

public class WCMapper extends MapReduceBase implements Mapper<LongWritable, Text, Text, IntWritable> {

    // Map function
    public void map(LongWritable key, Text value, OutputCollector<Text, IntWritable> output, Reporter reporter)
        throws IOException {
        String line = value.toString();

        // Splitting the line on whitespace
        for (String word : line.split("\\s+"))
        {
            if (word.length() > 0) {
                output.collect(new Text(word), new IntWritable(1));
            }
        }
    }
}

```

Reducer Code

```

import
java.io.IOException;
import java.util.Iterator;

import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import
org.apache.hadoop.mapred.MapReduceBase;
import org.apache.hadoop.mapred.OutputCollector;
import org.apache.hadoop.mapred.Reducer;
import org.apache.hadoop.mapred.Reporter;

public class WCReducer extends MapReduceBase implements Reducer<Text, IntWritable, Text, IntWritable> {

    // Reduce function
    public void reduce(Text key, Iterator<IntWritable>
        values, OutputCollector<Text, IntWritable>
        output, Reporter reporter) throws
        IOException {
        int count = 0;

        // Counting the frequency of each
        word while (values.hasNext()) {
            count += values.next().get();
        }

        output.collect(key, new IntWritable(count));
    }
}

```



```

hadoop@bmscsece-HP-Elite-Tower-600-G9-Desktop-PC:~$ hdfs WordCount
cd WordCount
hadoop@bmscsece-HP-Elite-Tower-600-G9-Desktop-PC:~/WordCount$ nano WMapper.java
hadoop@bmscsece-HP-Elite-Tower-600-G9-Desktop-PC:~/WordCount$ nano WReducer.java
hadoop@bmscsece-HP-Elite-Tower-600-G9-Desktop-PC:~/WordCount$ nano WDriver.java
hadoop@bmscsece-HP-Elite-Tower-600-G9-Desktop-PC:~/WordCount$ export HADOOP_CLASSPATH=$(hadoop classpath)
hadoop@bmscsece-HP-Elite-Tower-600-G9-Desktop-PC:~/WordCount$ javac -classpath $HADOOP_CLASSPATH -d . WMapper.java WReducer.java WDriver.java
hadoop@bmscsece-HP-Elite-Tower-600-G9-Desktop-PC:~/WordCount$ jar -cvf wordcount.jar *.class
added manifest
adding: WDriver.class(in = 1397) (out= 609)(deflated 56%)
adding: WMapper.class(in = 1857) (out= 757)(deflated 59%)
adding: WReducer.class(in = 1529) (out= 605)(deflated 60%)
hadoop@bmscsece-HP-Elite-Tower-600-G9-Desktop-PC:~/WordCount$ cd /usr/local/hadoop/sbin
start-all.sh
jps
bash: cd: /usr/local/hadoop/sbin: No such file or directory
WARNING: Attempting to start all Apache Hadoop daemons as hadoop in 10 seconds.
WARNING: This is not a recommended production deployment configuration.
WARNING: Use CTRL-C to abort.
Starting namenodes on [localhost]
Starting datanodes
Starting secondary namenodes [bmscsece-HP-Elite-Tower-600-G9-Desktop-PC]
Starting resourcemanager
Starting nodemanagers
9152 jps
8274 SecondaryNameNode
7977 DataNode
7882 NameNode
8576 ResourceManager
8732 NodeManager
hadoop@bmscsece-HP-Elite-Tower-600-G9-Desktop-PC:~/WordCount$ hdfs Input
nano Input/sample.txt
hadoop@bmscsece-HP-Elite-Tower-600-G9-Desktop-PC:~/WordCount$ hadoop fs -hdfs /Inputbda
hadoop fs -get Input/sample.txt /Inputbda
hadoop@bmscsece-HP-Elite-Tower-600-G9-Desktop-PC:~/WordCount$ hadoop jar wordcount.jar WDriver /Inputbda/sample.txt /Outputbda
2020-04-24 14:17:12,433 INFO Input.MetricsConfig: Loaded properties from hadoop-metrics2.properties
2020-04-24 14:17:12,478 INFO Input.MetricsSystemImpl: Scheduled metric snapshot period at 10 second(s).
2020-04-24 14:17:12,478 INFO Input.MetricsSystemImpl: JobTracker metrics system started
2020-04-24 14:17:12,477 WARN Input.MetricsSystemImpl: JobTracker metrics system already initialized!
2020-04-24 14:17:12,533 WARN Input.MetricsSystemImpl: Hadoop command-line option parsing not performed. Implement the Tool interface and execute your application with ToolRunner to remedy this.
2020-04-24 14:17:12,579 INFO Input.MetricsSystemImpl: Total Input Files to process : 1
2020-04-24 14:17:12,607 INFO Input.MetricsSystemImpl: number of splits:1
2020-04-24 14:17:12,646 INFO Input.MetricsSystemImpl: Submitting tokens for job: job_local138778888_0001
2020-04-24 14:17:12,646 INFO Input.MetricsSystemImpl: Submitting tokens for job: job_local138778888_0001
2020-04-24 14:17:12,720 INFO Input.MetricsSystemImpl: The url to track the job: http://localhost:8080/
2020-04-24 14:17:12,720 INFO Input.MetricsSystemImpl: Running job: job_local138778888_0001
2020-04-24 14:17:12,721 INFO Input.MetricsSystemImpl: OutputCommitter set in config null
2020-04-24 14:17:12,722 INFO Input.MetricsSystemImpl: OutputCommitter is org.apache.hadoop.mapred.FileOutputCommitter
2020-04-24 14:17:12,725 INFO Input.MetricsSystemImpl: File OutputCommitter Algorithm version is 2
2020-04-24 14:17:12,725 INFO Input.MetricsSystemImpl: skip cleanup _temporary folders under output directory:false, ignore cleanup failures: false
2020-04-24 14:17:12,762 INFO Input.MetricsSystemImpl: Waiting for map tasks
hadoop@bmscsece-HP-Elite-Tower-600-G9-Desktop-PC:~/WordCount$ hadoop fs -hdfs /Inputbda
hadoop fs -put Input/sample.txt /Inputbda
hadoop@bmscsece-HP-Elite-Tower-600-G9-Desktop-PC:~/WordCount$ hadoop jar wordcount.jar WDriver /Inputbda/sample.txt /Outputbda
2020-04-24 14:17:12,433 INFO Input.MetricsConfig: Loaded properties from hadoop-metrics2.properties
2020-04-24 14:17:12,478 INFO Input.MetricsSystemImpl: Scheduled metric snapshot period at 10 second(s).
2020-04-24 14:17:12,478 INFO Input.MetricsSystemImpl: JobTracker metrics system started
2020-04-24 14:17:12,477 WARN Input.MetricsSystemImpl: JobTracker metrics system already initialized!
2020-04-24 14:17:12,533 WARN Input.MetricsSystemImpl: Hadoop command-line option parsing not performed. Implement the Tool interface and execute your application with ToolRunner to remedy this.
2020-04-24 14:17:12,579 INFO Input.MetricsSystemImpl: Total Input Files to process : 1
2020-04-24 14:17:12,607 INFO Input.MetricsSystemImpl: number of splits:1
2020-04-24 14:17:12,646 INFO Input.MetricsSystemImpl: Submitting tokens for job: job_local138778888_0001
2020-04-24 14:17:12,646 INFO Input.MetricsSystemImpl: Submitting tokens for job: job_local138778888_0001
2020-04-24 14:17:12,720 INFO Input.MetricsSystemImpl: The url to track the job: http://localhost:8080/
2020-04-24 14:17:12,720 INFO Input.MetricsSystemImpl: Running job: job_local138778888_0001
2020-04-24 14:17:12,721 INFO Input.MetricsSystemImpl: OutputCommitter set in config null
2020-04-24 14:17:12,722 INFO Input.MetricsSystemImpl: OutputCommitter is org.apache.hadoop.mapred.FileOutputCommitter
2020-04-24 14:17:12,725 INFO Input.MetricsSystemImpl: File OutputCommitter Algorithm version is 2
2020-04-24 14:17:12,762 INFO Input.MetricsSystemImpl: skip cleanup _temporary folders under output directory:false, ignore cleanup failures: false
2020-04-24 14:17:12,764 INFO Input.MetricsSystemImpl: Starting task: attempt_local138778888_0001_m_000000_0
2020-04-24 14:17:12,775 INFO Input.MetricsSystemImpl: File OutputCommitter Algorithm version is 2
2020-04-24 14:17:12,775 INFO Input.MetricsSystemImpl: skip cleanup _temporary folders under output directory:false, ignore cleanup failures: false
2020-04-24 14:17:12,798 INFO Input.MetricsSystemImpl: Using ResourceCalculatorProcessTree : [ ]
2020-04-24 14:17:12,794 INFO Input.MetricsSystemImpl: Processing splits: hdfs://localhost:8080/Inputbda/sample.txt:0-49
2020-04-24 14:17:12,793 INFO Input.MetricsSystemImpl: numReduceTasks: 1
2020-04-24 14:17:12,821 INFO Input.MetricsSystemImpl: (0004708) @ kv1 26214396(104857184)
2020-04-24 14:17:12,821 INFO Input.MetricsSystemImpl: MapReduce task: kv1 180
2020-04-24 14:17:12,821 INFO Input.MetricsSystemImpl: kv1 limit at 83888888
2020-04-24 14:17:12,821 INFO Input.MetricsSystemImpl: bufstart = 0; bufend = 104857184
2020-04-24 14:17:12,821 INFO Input.MetricsSystemImpl: kvstart = 26214396; length = 83888888
2020-04-24 14:17:12,823 INFO Input.MetricsSystemImpl: Map output collector class = org.apache.hadoop.mapred.MapTask$MapOutputBuffer
2020-04-24 14:17:12,884 INFO Input.MetricsSystemImpl: Starting flush of map output
2020-04-24 14:17:12,884 INFO Input.MetricsSystemImpl: Spilling map output
2020-04-24 14:17:12,884 INFO Input.MetricsSystemImpl: bufstart = 0; bufend = 189; bufvold = 104857184
2020-04-24 14:17:12,884 INFO Input.MetricsSystemImpl: kvstart = 26214396(104857184); kvend = 26214396(104857184); length = 77/83888888
2020-04-24 14:17:12,887 INFO Input.MetricsSystemImpl: kv1 limit at 83888888
2020-04-24 14:17:12,892 INFO Input.MetricsSystemImpl: Task:attempt_local138778888_0001_m_000000_0 is done. And is in the process of committing
2020-04-24 14:17:12,894 INFO Input.MetricsSystemImpl: Task:attempt_local138778888_0001_m_000000_0 done.
2020-04-24 14:17:12,897 INFO Input.MetricsSystemImpl: Final Counters for attempt_local138778888_0001_m_000000_0: Counters: 23
File System Counters
FILE: Number of bytes read=2320
FILE: Number of bytes written=64187
FILE: Number of read operations=0
FILE: Number of large read operations=0
FILE: Number of write operations=0
HDFS: Number of bytes read=0
HDFS: Number of bytes written=0
HDFS: Number of read operations=0
HDFS: Number of large read operations=0
HDFS: Number of write operations=0
HDFS: Number of bytes read erasure-coded=0
Map-Reduce Framework
Map Input records=5

```



```

        Map input records=5
        Map output records=20
        Map output bytes=169
        Map output materialized bytes=215
        Input split bytes=93
        Combine input records=0
        Spilled Records=20
        Failed Shuffles=0
        Merged Map outputs=0
        GC time elapsed (ms)=0
        Total committed heap usage (bytes)=526385152
    File Input Format Counters
        Bytes Read=89
2026-04-24 14:17:12,897 INFO mapred.LocalJobRunner: Finishing task: attempt_local1387788888_0001_m_000000_0
2026-04-24 14:17:12,897 INFO mapred.LocalJobRunner: map task executor complete.
2026-04-24 14:17:12,899 INFO mapred.LocalJobRunner: Waiting for reduce tasks
2026-04-24 14:17:12,899 INFO mapred.LocalJobRunner: Starting task: attempt_local1387788888_0001_r_000000_0
2026-04-24 14:17:12,902 INFO output.FileOutputCommitter: File Output Committer Algorithm version is 2
2026-04-24 14:17:12,902 INFO output.FileOutputCommitter: FileOutputCommitter skip cleanup_temporary folders under output directory:false, ignore cleanup failures: false
2026-04-24 14:17:12,902 INFO mapred.Task: Using ResourceCalculatorProcessFree : [ ]
2026-04-24 14:17:12,904 INFO mapred.ReduceTask: Using ShuffleConsumerPlugin: org.apache.hadoop.mapreduce.task.reduce.Shuffle$53ac6b1c
2026-04-24 14:17:12,904 WARN Impl.MetricsSystemImpl: JobTracker metrics system already initialized!
2026-04-24 14:17:12,911 INFO reduce.MergeManagerImpl: MergeManager: memoryLimit=5830921216, maxSingleShuffleLimit=1457730304, mergeThreshold=3048400064, ioSortFactor=10, memToMemMergeOutputsThreshold=0
2026-04-24 14:17:12,912 INFO reduce.EventFetcher: attempt_local1387788888_0001_r_000000_0 Thread started: EventFetcher for fetching Map completion Events
2026-04-24 14:17:12,925 INFO reduce.LocalFetcher: localFetcher#1 about to shuffle output of map attempt_local1387788888_0001_m_000000_0 decomp: 211 len: 215 to MEMORY
2026-04-24 14:17:12,926 INFO reduce.InMemoryMapOutput: Read 211 bytes from map-output for attempt_local1387788888_0001_m_000000_0
2026-04-24 14:17:12,927 INFO reduce.MergeManagerImpl: closeInMemoryFile -> map-output of size: 211, inMemoryMapOutputs.size() -> 1, commitMemory -> 0, usedMemory -> 211
2026-04-24 14:17:12,927 INFO reduce.EventFetcher: EventFetcher is interrupted.. Returning
2026-04-24 14:17:12,928 INFO mapred.LocalJobRunner: 1 / 1 copied.
2026-04-24 14:17:12,928 INFO reduce.MergeManagerImpl: finalMerge called with 1 in-memory map-outputs and 0 on-disk map-outputs
2026-04-24 14:17:12,930 INFO mapred.Merger: Merging 1 sorted segments
2026-04-24 14:17:12,930 INFO mapred.Merger: Down to the last merge-pass, with 1 segments left of total size: 205 bytes
2026-04-24 14:17:12,931 INFO reduce.MergeManagerImpl: Merged 1 segments, 211 bytes to disk to satisfy reduce memory limit
2026-04-24 14:17:12,931 INFO reduce.MergeManagerImpl: Merging 1 files, 215 bytes from disk
2026-04-24 14:17:12,931 INFO reduce.MergeManagerImpl: Merging 0 segments, 0 bytes from memory into reduce
2026-04-24 14:17:12,931 INFO mapred.Merger: Merging 1 sorted segments
2026-04-24 14:17:12,931 INFO mapred.Merger: Down to the last merge-pass, with 1 segments left of total size: 205 bytes
2026-04-24 14:17:12,931 INFO mapred.LocalJobRunner: 1 / 1 copied.
2026-04-24 14:17:12,991 INFO mapred.Task: Task:attempt_local1387788888_0001_r_000000_0 is done. And is in the process of committing
2026-04-24 14:17:12,992 INFO mapred.LocalJobRunner: 1 / 1 copied.
2026-04-24 14:17:12,992 INFO mapred.Task: Task attempt_local1387788888_0001_r_000000_0 is allowed to commit now
2026-04-24 14:17:13,007 INFO output.FileOutputCommitter: Saved output of Task 'attempt_local1387788888_0001_r_000000_0' to hdfs://localhost:9000/outputbda
2026-04-24 14:17:13,008 INFO mapred.LocalJobRunner: reduce > reduce
2026-04-24 14:17:13,008 INFO mapred.Task: Task 'attempt_local1387788888_0001_r_000000_0' done.
2026-04-24 14:17:13,008 INFO mapred.Task: Final Counters for attempt_local1387788888_0001_r_000000_0: Counters: 30
    File System Counters
        FILE: Number of bytes read=3382
        FILE: Number of bytes written=644602
        FILE: Number of read operations=0
        FILE: Number of large read operations=0
        FILE: Number of write operations=0
        HDFS: Number of bytes read=89
        HDFS: Number of bytes written=69
        HDFS: Number of read operations=10
        HDFS: Number of large read operations=0
        HDFS: Number of write operations=3
        HDFS: Number of bytes read erasure-coded=0

```

```

2026-04-24 14:17:12,992 INFO mapred.LocalJobRunner: 1 / 1 copied.
2026-04-24 14:17:12,992 INFO mapred.Task: Task attempt_local1387788888_0001_r_000000_0 is allowed to commit now
2026-04-24 14:17:13,007 INFO output.FileOutputCommitter: Saved output of task 'attempt_local1387788888_0001_r_000000_0' to hdfs://localhost:9000/outputbda
2026-04-24 14:17:13,008 INFO mapred.LocalJobRunner: reduce > reduce
2026-04-24 14:17:13,008 INFO mapred.Task: Task 'attempt_local1387788888_0001_r_000000_0' done.
2026-04-24 14:17:13,008 INFO mapred.Task: Final Counters for attempt_local1387788888_0001_r_000000_0: Counters: 30
    File System Counters
        FILE: Number of bytes read=3382
        FILE: Number of bytes written=644602
        FILE: Number of read operations=0
        FILE: Number of large read operations=0
        FILE: Number of write operations=0
        HDFS: Number of bytes read=89
        HDFS: Number of bytes written=69
        HDFS: Number of read operations=10
        HDFS: Number of large read operations=0
        HDFS: Number of write operations=3
        HDFS: Number of bytes read erasure-coded=0

```

```

    File System Counters
        FILE: Number of bytes read=3382
        FILE: Number of bytes written=644602
        FILE: Number of read operations=0
        FILE: Number of large read operations=0
        FILE: Number of write operations=0
        HDFS: Number of bytes read=89
        HDFS: Number of bytes written=69
        HDFS: Number of read operations=10
        HDFS: Number of large read operations=0
        HDFS: Number of write operations=3
        HDFS: Number of bytes read erasure-coded=0
    Job Framework
        Combine input records=0
        Combine output records=0
        Reduce input groups=10
        Reduce shuffle bytes=215
        Reduce input records=20
        Reduce output records=10
        Spilled Records=20
        Shuffled Maps=1
        Failed Shuffles=0
        Merged Map outputs=1
        GC time elapsed (ms)=0
        Total committed heap usage (bytes)=526385152
    Shuffle Errors
        BAD_ID=0
        CONNECTION=0
        IO_ERROR=0
        WRONG_LENGTH=0
        WRONG_MAP=0
        WRONG_REDUCE=0
    File Output Format Counters
        Bytes Written=69
2026-04-24 14:17:13,008 INFO mapred.LocalJobRunner: Finishing task: attempt_local1387788888_0001_r_000000_0
2026-04-24 14:17:13,008 INFO mapred.LocalJobRunner: reduce task executor complete.
2026-04-24 14:17:13,723 INFO mapreduce.Job: Job job_local1387788888_0001 running in uber mode : false
2026-04-24 14:17:13,723 INFO mapreduce.Job: map 100% reduce 100%
2026-04-24 14:17:13,724 INFO mapreduce.Job: Job job_local1387788888_0001 completed successfully
2026-04-24 14:17:13,720 INFO mapreduce.Job: Counters: 36
    File System Counters
        FILE: Number of bytes read=6302
        FILE: Number of bytes written=1288989
        FILE: Number of read operations=0
        FILE: Number of large read operations=0
        FILE: Number of write operations=0
        HDFS: Number of bytes read=178
        HDFS: Number of bytes written=69
        HDFS: Number of read operations=15

```

```

FILE: Number of bytes written=1288989
FILE: Number of read operations=0
FILE: Number of large read operations=0
FILE: Number of write operations=0
HDFS: Number of bytes read=178
HDFS: Number of bytes written=69
HDFS: Number of read operations=15
HDFS: Number of large read operations=0
HDFS: Number of write operations=4
HDFS: Number of bytes read erasure-coded=0
Map-Reduce Framework
  Map input records=5
  Map output records=20
  Map output bytes=169
  Map output materialized bytes=215
  Input split bytes=93
  Combine input records=0
  Combine output records=0
  Reduce input groups=10
  Reduce shuffle bytes=215
  Reduce input records=20
  Reduce output records=10
  Spilled Records=40
  Shuffled Maps =1
  Failed Shuffles=0
  Merged Map outputs=1
  GC time elapsed (ms)=0
  Total committed heap usage (bytes)=1052770304
Shuffle Errors
  BAD_ID=0
  CONNECTION=0
  IO_ERROR=0
  WRONG_LENGTH=0
  WRONG_MAP=0
  WRONG_REDUCE=0
File Input Format Counters
  Bytes Read=89
File Output Format Counters
  Bytes Written=69
hadoop@bmsccscse-HP-Elite-Tower-600-G9-Desktop-PC: ~/WordCount$ hadoop fs -ls /outputbda
Found 2 items
-rw-r--r-- 1 hadoop supergroup      0 2026-04-24 14:17 /outputbda/_SUCCESS
-rw-r--r-- 1 hadoop supergroup    69 2026-04-24 14:17 /outputbda/part-00000
hadoop@bmsccscse-HP-Elite-Tower-600-G9-Desktop-PC: ~/WordCount$ hadoop fs -cat /outputbda/part-00000
are 1
brother 1
family 1
hi 1
how 5
is 4
job 1
sister 1
you 1
your 4
hadoop@bmsccscse-HP-Elite-Tower-600-G9-Desktop-PC: ~/WordCount$

```

PROGRAM – 6

Implement Weather program on Hadoop framework

Questions:

From the following link extract the weather data

<https://github.com/tomwhite/hadoopbook/tree/master/input/ncdc/all>

- a) Create a MapReduce program to find average temperature for each year from NCDC data set.
- b) find the mean max temperature for every month.

Driver Code

```
package temp;

import
org.apache.hadoop.conf.Configuration;
import org.apache.hadoop.fs.Path;

import
org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;

import org.apache.hadoop.mapreduce.Job;

import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
import
org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;

public class AverageDriver {
    public static void main(String[] args) throws Exception {
        if (args.length != 2) {
            System.err.println("Please enter both input and output
            parameters."); System.exit(-1);
        }

        // Creating a configuration and job
        instance Configuration conf = new
        Configuration();
        Job job = Job.getInstance(conf, "Average

        Calculation"); job.setJarByClass(AverageDriver.class);

        // Input and output paths
        FileInputFormat.addInputPath(job, new Path(args[0]));
        FileOutputFormat.setOutputPath(job, new
        Path(args[1]));

        // Setting mapper and reducer classes
        job.setMapperClass(AverageMapper.class)
        ;
        job.setReducerClass(AverageReducer.clas
        s);

        // Output key and value types
        job.setOutputKeyClass(Text.class);
        job.setOutputValueClass(IntWritable.clas
        s);

        // Submitting the job and waiting for it to complete
```

```
        System.exit(job.waitForCompletion(true) ? 0 : 1);  
    }  
}
```

Mapper Code

```
package temp;

import java.io.IOException;

import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.LongWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Mapper;

public class AverageMapper extends Mapper<LongWritable, Text, Text, IntWritable>

{ public static final int MISSING = 9999;

    @Override
    public void map(LongWritable key, Text value, Context
        context) throws IOException, InterruptedException {

        String line = value.toString();

        // Extract year from fixed position
        String year = line.substring(15,
            19); int temperature;

        // Determine if there's a '+'
        sign if (line.charAt(87) ==
            '+') {
            temperature = Integer.parseInt(line.substring(88, 92));
        } else {
            temperature = Integer.parseInt(line.substring(87, 92));
        }

        // Quality check character
        String quality = line.substring(92, 93);

        // Only emit if data is valid
        if (temperature != MISSING && quality.matches("[01459]")) {
            context.write(new Text(year), new IntWritable(temperature));
        }
    }
}
```

Reducer Code

```
package temp;

import java.io.IOException;

import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Reducer;

public class AverageReducer extends Reducer<Text, IntWritable, Text, IntWritable> {
    @Override
    public void reduce(Text key, Iterable<IntWritable> values,
        Context context) throws IOException, InterruptedException {

        int sumTemp =
            0; int count = 0;

        for (IntWritable value : values)
            { sumTemp += value.get();
              count++;
            }

        if (count > 0) {
            int average = sumTemp / count;
            context.write(key, new IntWritable(average));
        }
    }
}
```


Driver Code

```
package meanmax;

import
org.apache.hadoop.conf.Configuration;
import org.apache.hadoop.fs.Path;

import
org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;

import org.apache.hadoop.mapreduce.Job;

import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
import
org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;

public class MeanMaxDriver {
    public static void main(String[] args) throws Exception {
        if (args.length != 2) {
            System.err.println("Please enter both input and output
            parameters."); System.exit(-1);
        }

        Configuration conf = new Configuration();
        Job job = Job.getInstance(conf, "Mean and Max
        Temperature"); job.setJarByClass(MeanMaxDriver.class);

        FileInputFormat.addInputPath(job, new Path(args[0]));
        FileOutputFormat.setOutputPath(job, new
        Path(args[1]));

        job.setMapperClass(MeanMaxMapper.class
        );
        job.setReducerClass(MeanMaxReducer.clas
        s);

        job.setOutputKeyClass(Text.class);
        job.setOutputValueClass(IntWritable.clas
        s);

        System.exit(job.waitForCompletion(true) ? 0 : 1);
    }
}
```

Mapper Code

```
package meanmax;

import java.io.IOException;

import org.apache.hadoop.io.IntWritable;
import
org.apache.hadoop.io.LongWritable;
import org.apache.hadoop.io.Text;

import org.apache.hadoop.mapreduce.Mapper;

public class MeanMaxMapper extends Mapper<LongWritable, Text, Text, IntWritable>

{ public static final int MISSING = 9999;

@Override
public void map(LongWritable key, Text value, Context
context) throws IOException, InterruptedException {

    String line = value.toString();
```



```
// Extract month from positions 19-20
String month = line.substring(19, 21);
int temperature;
```

```

// Extract temperature considering optional
'+' if (line.charAt(87) == '+') {
    temperature = Integer.parseInt(line.substring(88, 92));
} else {
    temperature = Integer.parseInt(line.substring(87, 92));
}

// Quality check
String quality = line.substring(92, 93);

if (temperature != MISSING && quality.matches("[01459]")) {
    context.write(new Text(month), new
        IntWritable(temperature));
}
}
}

```

Reducer Code

```

package meanmax;

import java.io.IOException;

import
org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Reducer;
public class MeanMaxReducer extends Reducer<Text, IntWritable, Text, Text> {
    @Override
    public void reduce(Text key, Iterable<IntWritable> values,
        Context context) throws IOException, InterruptedException {

        int sumTemp =
        0; int count = 0;
        int maxTemp = Integer.MIN_VALUE;

        for (IntWritable value : values)
        { int temp = value.get();
          sumTemp += temp;
          count++;

          if (temp > maxTemp)
            { maxTemp = temp;
            }
        }

        if (count > 0) {
            int avgTemp = sumTemp / count;
            String result = "mean=" + avgTemp + " max=" + maxTemp;
            context.write(key, new Text(result));
        }
    }
}

```

```

1901 46
1902 21
hadoop@bmsccse-HP-Elite-Tower-600-G9-Desktop-PC: $ hdfs dfs -rm -r /meannaxoutput
rm: '/meannaxoutput': No such file or directory
hadoop@bmsccse-HP-Elite-Tower-600-G9-Desktop-PC: $ hadoop jar Weather.jar meannax.MeanMaxDriver /weather /meannaxoutput
2026-03-13 16:05:30,143 INFO client.DefaultHARMFallOverProxyProvider: Connecting to ResourceManager at /0.0.0.0:8032
2026-03-13 16:05:30,276 WARN mapreduce.JobResourceUploader: Hadoop command-line option parsing not performed. Implement the Tool interface and execute your application with ToolRunner to remedy this.
2026-03-13 16:05:30,305 INFO mapreduce.JobResourceUploader: Disabling Erasure Coding for path: /tmp/hadoop-yarn/staging/hadoop/.staging/job_1773395565039_0002
2026-03-13 16:05:30,453 INFO input.FileInputFormat: Total input files to process : 2
2026-03-13 16:05:30,529 INFO mapreduce.JobSubmitter: Number of splits:2
2026-03-13 16:05:30,612 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1773395565039_0002
2026-03-13 16:05:30,612 INFO mapreduce.JobSubmitter: Executing with tokens: []
2026-03-13 16:05:30,682 INFO conf.Configuration: resource-types.xml not found
2026-03-13 16:05:30,682 INFO resource.ResourceUtils: Unable to find 'resource-types.xml'.
2026-03-13 16:05:30,709 INFO impl.VarClientImpl: Submitted application application_1773395565039_0002
2026-03-13 16:05:30,725 INFO mapreduce.Job: The url to track the job: http://bmsccse-HP-Elite-Tower-600-G9-Desktop-PC:8088/proxy/application_1773395565039_0002/
2026-03-13 16:05:30,726 INFO mapreduce.Job: Running job: job_1773395565039_0002
2026-03-13 16:05:34,760 INFO mapreduce.Job: Job job_1773395565039_0002 running in uber mode : false
2026-03-13 16:05:34,761 INFO mapreduce.Job: map 0% reduce 0%
2026-03-13 16:05:37,786 INFO mapreduce.Job: map 100% reduce 0%
2026-03-13 16:05:41,800 INFO mapreduce.Job: map 100% reduce 100%
2026-03-13 16:05:41,806 INFO mapreduce.Job: Job job_1773395565039_0002 completed successfully
2026-03-13 16:05:41,847 INFO mapreduce.Job: Counters: 54

File System Counters
  FILE: Number of bytes read=118167
  FILE: Number of bytes written=1064069
  FILE: Number of read operations=0
  FILE: Number of large read operations=0
  FILE: Number of write operations=0
  HDFS: Number of bytes read=1773366
  HDFS: Number of bytes written=72
  HDFS: Number of read operations=11
  HDFS: Number of large read operations=0
  HDFS: Number of write operations=2
  HDFS: Number of bytes read erasure-coded=0

Job Counters
  Launched map tasks=2
  Launched reduce tasks=1
  Data-local map tasks=2
  Total time spent by all maps in occupied slots (ms)=2213
  Total time spent by all reduces in occupied slots (ms)=1047
  Total time spent by all map tasks (ms)=2213
  Total time spent by all reduce tasks (ms)=1047
  Total vcore-millisecods taken by all map tasks=2213
  Total vcore-millisecods taken by all reduce tasks=1047
  Total megabyte-millisecods taken by all map tasks=2266112
  Total megabyte-millisecods taken by all reduce tasks=1072128

Map-Reduce Framework
  Map input records=13130
  Map output records=13129
  Map output bytes=91903
  Map output materialized bytes=118173
  Input split bytes=198
  Combine input records=0
  Combine output records=0

Total time spent by all reduce tasks (ms)=1047
Total vcore-millisecods taken by all map tasks=2213
Total vcore-millisecods taken by all reduce tasks=1047
Total megabyte-millisecods taken by all map tasks=2266112
Total megabyte-millisecods taken by all reduce tasks=1072128

Map-Reduce Framework
  Map input records=13130
  Map output records=13129
  Map output bytes=91903
  Map output materialized bytes=118173
  Input split bytes=198
  Combine input records=0
  Combine output records=0
  Reduce input groups=12
  Reduce shuffle bytes=118173
  Reduce input records=13129
  Reduce output records=12
  Spilled Records=26258
  Shuffled Maps =2
  Failed Shuffles=0
  Merged Map outputs=2
  GC time elapsed (ms)=45
  CPU time spent (ms)=2030
  Physical memory (bytes) snapshot=1037684736
  Virtual memory (bytes) snapshot=8376963072
  Total committed heap usage (bytes)=1572864000
  Peak Map Physical memory (bytes)=381546496
  Peak Map Virtual memory (bytes)=2792185856
  Peak Reduce Physical memory (bytes)=281702400
  Peak Reduce Virtual memory (bytes)=2799857664

Shuffle Errors
  BAD_ID=0
  CONNECTION=0
  IO_ERROR=0
  WRONG_LENGTH=0
  WRONG_MAP=0
  WRONG_REDUCE=0

File Input Format Counters
  Bytes Read=1773366
File Output Format Counters
  Bytes Written=72
hadoop@bmsccse-HP-Elite-Tower-600-G9-Desktop-PC: $ hdfs dfs -cat /meannaxoutput/part-r-00000
01 4
02 1
03 6
04 34
05 89
06 143
07 182
08 172
09 123
10 73
11 21
12 3
hadoop@bmsccse-HP-Elite-Tower-600-G9-Desktop-PC: $

```

```

Total time spent by all reduce tasks (ms)=1047
Total vcore-millisecods taken by all map tasks=2213
Total vcore-millisecods taken by all reduce tasks=1047
Total megabyte-millisecods taken by all map tasks=2266112
Total megabyte-millisecods taken by all reduce tasks=1072128

Map-Reduce Framework
  Map input records=13130
  Map output records=13129
  Map output bytes=91903
  Map output materialized bytes=118173
  Input split bytes=198
  Combine input records=0
  Combine output records=0
  Reduce input groups=12
  Reduce shuffle bytes=118173
  Reduce input records=13129
  Reduce output records=12
  Spilled Records=26258
  Shuffled Maps =2
  Failed Shuffles=0
  Merged Map outputs=2
  GC time elapsed (ms)=45
  CPU time spent (ms)=2030
  Physical memory (bytes) snapshot=1037684736
  Virtual memory (bytes) snapshot=8376963072
  Total committed heap usage (bytes)=1572864000
  Peak Map Physical memory (bytes)=381546496
  Peak Map Virtual memory (bytes)=2792185856
  Peak Reduce Physical memory (bytes)=281702400
  Peak Reduce Virtual memory (bytes)=2799857664

Shuffle Errors
  BAD_ID=0
  CONNECTION=0
  IO_ERROR=0
  WRONG_LENGTH=0
  WRONG_MAP=0
  WRONG_REDUCE=0

File Input Format Counters
  Bytes Read=1773366
File Output Format Counters
  Bytes Written=72
hadoop@bmsccse-HP-Elite-Tower-600-G9-Desktop-PC: $ hdfs dfs -cat /meannaxoutput/part-r-00000
01 4
02 1
03 6
04 34
05 89
06 143
07 182
08 172
09 123
10 73
11 21
12 3
hadoop@bmsccse-HP-Elite-Tower-600-G9-Desktop-PC: $

```

PROGRAM – 7

For a given Text file, Create a Map Reduce program to sort the content in an alphabetic order listing only top 10 maximum occurrences of words.

Driver Code (TopNDriver.java)

```
package samples.topn;

import
org.apache.hadoop.conf.Configuration;
import org.apache.hadoop.fs.Path;

import
org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;

import org.apache.hadoop.mapreduce.Job;
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
import
org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;

public class TopNDriver {

    public static void main(String[] args) throws Exception {
        if (args.length != 3) {
            System.err.println("Usage: TopNDriver <in> <temp-out>
<final-out>"); System.exit(2);
        }

        Configuration conf = new Configuration();

        // === Job 1: Word Count ===
        Job wcJob = Job.getInstance(conf, "word count");
        wcJob.setJarByClass(TopNDriver.class);
        wcJob.setMapperClass(WordCountMapper.class);
        wcJob.setCombinerClass(WordCountReducer.class);
        wcJob.setReducerClass(WordCountReducer.class);
        wcJob.setOutputKeyClass(Text.class);
        wcJob.setOutputValueClass(IntWritable.class);

        FileInputFormat.addInputPath(wcJob, new Path(args[0]));
        Path tempDir = new Path(args[1]);
        FileOutputFormat.setOutputPath(wcJob, tempDir);

        if (!wcJob.waitForCompletion(true))
            { System.exit(1);
            }

        // === Job 2: Top N ===
        Job topJob = Job.getInstance(conf, "top 10
words"); topJob.setJarByClass(TopNDriver.class);
        topJob.setMapperClass(TopNMapper.class);
        topJob.setReducerClass(TopNReducer.class);
        topJob.setMapOutputKeyClass(IntWritable.class);
        topJob.setMapOutputValueClass(Text.class);
        topJob.setOutputKeyClass(Text.class);
        topJob.setOutputValueClass(IntWritable.class);

        FileInputFormat.addInputPath(topJob, tempDir);
        FileOutputFormat.setOutputPath(topJob, new
Path(args[2]));
    }
}
```

```
        System.exit(topJob.waitForCompletion(true) ? 0 : 1);  
    }  
}
```

Mapper Code (WordCountMapper.java)

```
package samples.topn;

import java.io.IOException;
import java.util.StringTokenizer;

import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Mapper;

public class WordCountMapper
    extends Mapper<Object, Text, Text, IntWritable> {

    private final static IntWritable ONE = new
        IntWritable(1); private Text word = new Text();
    // characters to normalize into spaces
    private String tokens = "[_!$%&'\\";

    @Override
    protected void map(Object key, Text value, Context
        context) throws IOException, InterruptedException {

        // clean & tokenize
        String clean = value.toString()
            .toLowerCase()
            .replaceAll(tokens, " ");
        StringTokenizer itr = new
            StringTokenizer(clean); while
            (itr.hasMoreTokens()) {
                word.set(itr.nextToken().trim());
                context.write(word, ONE);
            }
    }
}
```

Mapper Code (TopNMapper.java)

```
package samples.topn;

import java.io.IOException;

import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Mapper;

public class TopNMapper
    extends Mapper<Object, Text, IntWritable, Text> {

    private IntWritable count = new
        IntWritable(); private Text word = new
        Text();

    @Override
    protected void map(Object key, Text value, Context
        context) throws IOException, InterruptedException {

        // input line: word \t count
        String[] parts =
            value.toString().split("\\t"); if (parts.length
            == 2) {
                word.set(parts[0]);
                count.set(Integer.parseInt(parts[1])
                );
                // emit count → word, so Hadoop sorts by count
                context.write(count, word);
            }
    }
}
```

```
}
```

Reducer Code (WordCountReducer.java)

```
package samples.topn;
```

```
import
```

```
java.io.IOException;
```



```

import
org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Reducer;

public class WordCountReducer
    extends Reducer<Text, IntWritable, Text, IntWritable> {

    @Override
    protected void reduce(Text key, Iterable<IntWritable> values, Context
        context) throws IOException, InterruptedException {

        int sum = 0;
        for (IntWritable val : values)
            { sum += val.get();
            }
        context.write(key, new IntWritable(sum));
    }
}

```

Reducer Code (TopNReducer.java)

```

package samples.topn;

import
java.io.IOException;
import java.util.ArrayList;
import
java.util.Collections;
import java.util.List;
import java.util.Map;
import java.util.TreeMap;

import
org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Reducer;

public class TopNReducer
    extends Reducer<IntWritable, Text, Text, IntWritable> {

    // TreeMap with descending order of keys (counts)
    private TreeMap<Integer, List<String>> countMap
    =
        new TreeMap<>(Collections.reverseOrder());

    @Override
    protected void reduce(IntWritable key, Iterable<Text> values, Context
        context) throws IOException, InterruptedException {

        int cnt = key.get();
        List<String> words = countMap.getOrDefault(cnt, new
            ArrayList<>()); for (Text w : values) {
            words.add(w.toString());
        }
        countMap.put(cnt, words);
    }

    @Override
    protected void cleanup(Context context)
        throws IOException, InterruptedException
    {

        // collect top 10 word→count pairs
        List<WordCount> topList = new
            ArrayList<>(); int seen = 0;
        for (Map.Entry<Integer, List<String>> entry : countMap.entrySet())
            { int cnt = entry.getKey();
            for (String w : entry.getValue()) {
                topList.add(new WordCount(w,
                    cnt)); seen++;
                if (seen == 10) break;
            }
        }
    }
}

```

```
}  
if (seen == 10) break;
```

```

    }

    // sort these 10 entries alphabetically by word
    Collections.sort(topList, (a, b) -> a.word.compareTo(b.word));

    // emit final top 10 in alphabetical
    order for (WordCount wc : topList) {
        context.write(new Text(wc.word), new IntWritable(wc.count));
    }
}

// helper class
private static class WordCount
{
    String word;
    int count;
    WordCount(String w, int c) { word = w; count = c; }
}
}

```

The screenshot shows a terminal window with the GNU nano 6.2 text editor. The title bar at the top reads "GNU nano 6.2" and "/home/hadoop/sample.txt". The editor contains five lines of text: "hi how are you", "how is your job", "how is your family", "how is your brother", and "how is your sister". At the bottom of the editor, there is a status bar that says "[Read 5 lines]". Below the status bar is a toolbar with various keyboard shortcuts: ^G Help, ^O Write Out, ^W Where Is, ^K Cut, ^T Execute, ^C Location, ^X Exit, ^R Read File, ^\ Replace, ^U Paste, ^J Justify, and ^_ Go To Line.

```

GNU nano 6.2 /home/hadoop/sample.txt
hi how are you
how is your job
how is your family
how is your brother
how is your sister

[ Read 5 lines ]
^G Help    ^O Write Out  ^W Where Is  ^K Cut      ^T Execute   ^C Location
^X Exit    ^R Read File  ^\ Replace   ^U Paste    ^J Justify   ^_ Go To Line

```

```

hadoop@bmscscce-HP-Elite-Tower-600-G9-Desktop-PC: ~/hadoop/sbin$ ./start-all.sh
WARNING: Attempting to start all Apache Hadoop daemons as hadoop in 10 seconds.
WARNING: This is not a recommended production deployment configuration.
WARNING: Use CTRL-C to abort.
Starting namenodes on [localhost]
Starting datanodes
Starting secondary namenodes [bmscscce-HP-Elite-Tower-600-G9-Desktop-PC]
Starting resourcemanager
Starting nodemanagers
hadoop@bmscscce-HP-Elite-Tower-600-G9-Desktop-PC:~/hadoop/sbin$ jps
31057 Jps
30787 NodeManager
30149 SecondaryNameNode
29706 NameNode
29902 DataNode
30431 ResourceManager
hadoop@bmscscce-HP-Elite-Tower-600-G9-Desktop-PC:~/hadoop/sbin$ cd ~/topn
hadoop@bmscscce-HP-Elite-Tower-600-G9-Desktop-PC:~/topn$ javac -classpath 'hadoop classpath' -d . TopN.java TopMapper.java TopReducer.java TopCombiner.java utils/MiscUtils.java
Note: utils/MiscUtils.java uses unchecked or unsafe operations.
Note: Recompile with -Xlint:unchecked for details.
hadoop@bmscscce-HP-Elite-Tower-600-G9-Desktop-PC:~/topn$ jar -cvf TopN.jar *.class utils/*.class
added manifest
adding: TopN.class(in = 1308) (out= 764)(deflated 44%)
adding: TopCombiner.class(in = 1587) (out= 662)(deflated 58%)
adding: TopMapper.class(in = 1953) (out= 883)(deflated 54%)
adding: TopReducer.class(in = 2532) (out= 1003)(deflated 60%)
adding: utils/MiscUtils$1.class(in = 944) (out= 494)(deflated 47%)
adding: utils/MiscUtils.class(in = 1231) (out= 693)(deflated 43%)
hadoop@bmscscce-HP-Elite-Tower-600-G9-Desktop-PC:~/topn$ nano ~/sample.txt
mkdir: '/topnInput': File exists
hadoop@bmscscce-HP-Elite-Tower-600-G9-Desktop-PC:~/topn$ hadoop fs -put ~/sample.txt /topninput
put: '/topninput/sample.txt': File exists
hadoop@bmscscce-HP-Elite-Tower-600-G9-Desktop-PC:~/topn$ hadoop fs -ls /topninput
Found 1 items
-rw-r--r-- 1 hadoop supergroup      89 2026-04-24 14:58 /topninput/sample.txt
hadoop@bmscscce-HP-Elite-Tower-600-G9-Desktop-PC:~/topn$ hadoop jar TopN.jar TopN /topninput/sample.txt /topnoutput
2026-04-24 15:04:22,185 INFO client.DefaultHoharmFalloverProxyProvider: Connecting to ResourceManager at /0.0.0.0:8032
Exception in thread "main" org.apache.hadoop.mapred.FileAlreadyExistsException: Output directory hdfs://localhost:9000/topnoutput already exists
    at org.apache.hadoop.mapreduce.lib.output.FileOutputFormat.checkOutputSpecs(FileOutputFormat.java:164)
    at org.apache.hadoop.mapreduce.JobSubmitter.checkSpecs(JobSubmitter.java:277)
    at org.apache.hadoop.mapreduce.JobSubmitter.submitJobInternal(JobSubmitter.java:143)
    at org.apache.hadoop.mapreduce.Job$11.run(Job.java:1678)
    at org.apache.hadoop.mapreduce.Job$11.run(Job.java:1675)
    at java.base/java.security.AccessController.doPrivileged(Native Method)
    at java.base/javax.security.auth.Subject.doAs(Subject.java:423)
    at org.apache.hadoop.security.UserGroupInformation.doAs(UserGroupInformation.java:1899)
    at org.apache.hadoop.mapreduce.Job.submit(Job.java:1675)
    at org.apache.hadoop.mapreduce.Job.waitForCompletion(Job.java:1696)
    at TopN.main(TopN.java:27)
    at java.base/jdk.internal.reflect.NativeMethodAccessorImpl.invoke0(Native Method)

```

```

Note: utils/MiscUtils.java uses unchecked or unsafe operations.
Note: Recompile with -Xlint:unchecked for details.
hadoop@bmscscce-HP-Elite-Tower-600-G9-Desktop-PC:~/topn$ jar -cvf TopN.jar *.class utils/*.class
added manifest
adding: TopN.class(in = 1308) (out= 764)(deflated 44%)
adding: TopCombiner.class(in = 1587) (out= 662)(deflated 58%)
adding: TopMapper.class(in = 1953) (out= 883)(deflated 54%)
adding: TopReducer.class(in = 2532) (out= 1003)(deflated 60%)
adding: utils/MiscUtils$1.class(in = 944) (out= 494)(deflated 47%)
adding: utils/MiscUtils.class(in = 1231) (out= 693)(deflated 43%)
hadoop@bmscscce-HP-Elite-Tower-600-G9-Desktop-PC:~/topn$ nano ~/sample.txt
hadoop@bmscscce-HP-Elite-Tower-600-G9-Desktop-PC:~/topn$ hadoop fs -mkdir /topninput
mkdir: '/topnInput': File exists
hadoop@bmscscce-HP-Elite-Tower-600-G9-Desktop-PC:~/topn$ hadoop fs -put ~/sample.txt /topninput
put: '/topninput/sample.txt': File exists
hadoop@bmscscce-HP-Elite-Tower-600-G9-Desktop-PC:~/topn$ hadoop fs -ls /topninput
Found 1 items
-rw-r--r-- 1 hadoop supergroup      89 2026-04-24 14:58 /topninput/sample.txt
hadoop@bmscscce-HP-Elite-Tower-600-G9-Desktop-PC:~/topn$ hadoop jar TopN.jar TopN /topninput/sample.txt /topnoutput
2026-04-24 15:04:22,185 INFO client.DefaultHoharmFalloverProxyProvider: Connecting to ResourceManager at /0.0.0.0:8032
Exception in thread "main" org.apache.hadoop.mapred.FileAlreadyExistsException: Output directory hdfs://localhost:9000/topnoutput already exists
    at org.apache.hadoop.mapreduce.lib.output.FileOutputFormat.checkOutputSpecs(FileOutputFormat.java:164)
    at org.apache.hadoop.mapreduce.JobSubmitter.checkSpecs(JobSubmitter.java:277)
    at org.apache.hadoop.mapreduce.JobSubmitter.submitJobInternal(JobSubmitter.java:143)
    at org.apache.hadoop.mapreduce.Job$11.run(Job.java:1678)
    at org.apache.hadoop.mapreduce.Job$11.run(Job.java:1675)
    at java.base/java.security.AccessController.doPrivileged(Native Method)
    at java.base/javax.security.auth.Subject.doAs(Subject.java:423)
    at org.apache.hadoop.security.UserGroupInformation.doAs(UserGroupInformation.java:1899)
    at org.apache.hadoop.mapreduce.Job.submit(Job.java:1675)
    at org.apache.hadoop.mapreduce.Job.waitForCompletion(Job.java:1696)
    at TopN.main(TopN.java:27)
    at java.base/jdk.internal.reflect.NativeMethodAccessorImpl.invoke0(Native Method)
    at java.base/jdk.internal.reflect.NativeMethodAccessorImpl.invoke(NativeMethodAccessorImpl.java:62)
    at java.base/jdk.internal.reflect.DelegatingMethodAccessorImpl.invoke(DelegatingMethodAccessorImpl.java:43)
    at java.base/java.lang.reflect.Method.invoke(Method.java:566)
    at org.apache.hadoop.util.RunJar.run(RunJar.java:328)
    at org.apache.hadoop.util.RunJar.main(RunJar.java:241)
hadoop@bmscscce-HP-Elite-Tower-600-G9-Desktop-PC:~/topn$ hadoop fs -cat /topnoutput/part-r-00000
how      5
your    4
is       4
brother 1
are      1
hl       1
sister  1
family  1
you      1
job      1
hadoop@bmscscce-HP-Elite-Tower-600-G9-Desktop-PC:~/topn$

```

PROGRAM – 8

Write a Scala program to print numbers from 1 to 100 using for loop.

```
object ExampleForLoop1 {  
  def main(args: Array[String]) { for(counter  
    <- 1 to 100) {  
    print(counter + " ")  
    }  
    println()  
  }  
}
```

Output:

```
hadoop@bmscscse-HP-Elite-Tower-600-G9-Desktop-PC: $ scala -version  
Scala code runner version 2.12.18 -- Copyright 2002-2023, LAMP/EPFL and Lightbend, Inc.  
hadoop@bmscscse-HP-Elite-Tower-600-G9-Desktop-PC: $ nano PrintNumbers.scala  
hadoop@bmscscse-HP-Elite-Tower-600-G9-Desktop-PC: $ scala PrintNumbers.scala  
1  
2  
3  
4  
5  
6  
7  
8  
9  
10  
11  
12  
13  
14  
15  
16  
17  
18  
19  
20  
21  
22  
23  
24  
25  
26  
27  
28  
29  
30  
31  
32  
33  
34  
35  
36  
37  
38  
39  
40  
41  
42  
43  
44  
45  
46  
47
```

```
49  
50  
51  
52  
53  
54  
55  
56  
57  
58  
59  
60  
61  
62  
63  
64  
65  
66  
67  
68  
69  
70  
71  
72  
73  
74  
75  
76  
77  
78  
79  
80  
81  
82  
83  
84  
85  
86  
87  
88  
89  
90  
91  
92  
93  
94  
95  
96  
97  
98  
99  
100  
hadoop@bmscscse-HP-Elite-Tower-600-G9-Desktop-PC: $
```

PROGRAM – 9

Using RDD and FlatMap count how many times each word appears in a file and write out a list of words whose count is strictly greater than 4 using Spark.

```
scala> val textFile = sc.textFile("D:\\Downloads\\spark-4.1.1-bin-hadoop3\\spark-4.1.1-bin-hadoop3\\spark-data\\wc.txt")
val textFile: org.apache.spark.rdd.RDD[String] = D:\\Downloads\\spark-4.1.1-bin-hadoop3\\spark-4.1.1-bin-hadoop3\\spark-data\\wc.txt MapPartitionsRDD[1] at textF
ile at <console>:1

scala> val counts = textFile.flatMap(line => line.split(" ")).map(word => (word, 1)).reduceByKey(_ + _)
val counts: org.apache.spark.rdd.RDD[(String, Int)] = ShuffledRDD[4] at reduceByKey at <console>:1

scala> val sorted = counts.collect().sortBy(_._2)
val sorted: Array[(String, Int)] = Array((spark,6), (hello,4), (java,3), (python,2), (world,1))

scala> sorted.foreach(println)
(spark,6)
(hello,4)
(java,3)
(python,2)
(world,1)

scala> for((k,v) <- sorted) {
  |   if(v > 4) {
  |     print(k + " -> ")
  |     print(v)
  |     println()
  |   }
  | }
spark -> 6

scala> sorted.filter(_._2 > 4).foreach(println)
(spark,6)
```

PROGRAM – 10

Write a simple streaming program in Spark to receive text data streams on a particular port, perform basic text cleaning (like white space removal, stop words removal, lemmatization, etc.), and print the cleaned text on the screen. (Open Ended Question).

```
C:\Users\91988>spark-shell
WARNING: Using incubator modules: jdk.incubator.vector
Using Spark's default log4j profile: org/apache/spark/log4j2-defaults.properties
Setting default log level to "WARN".
To adjust logging level use sc.setLogLevel(newLevel). For SparkR, use setLogLevel(newLevel).
Welcome to

  ____      _
 / ___|    / \
| |  | |  / _ \
| |  | | / ___ \
| |  | |/_/   \_\
|_|  |_____/___\_\

version 4.1.1

Using Scala version 2.13.17 (OpenJDK 64-Bit Server VM, Java 17.0.18)
Type in expressions to have them evaluated.
Type :help for more information.
Spark context Web UI available at http://LAPTOP-6DEICVVQ:4848
Spark context available as 'sc' (master = local[*], app id = local-1777009385296).
Spark session available as 'spark'.

scala> import spark.implicits._
import spark.implicits._

scala> spark.sparkContext.setLogLevel("ERROR")

scala> val lines = spark.readStream.format("socket").option("host", "localhost").option("port", 9999).load()
val lines: org.apache.spark.sql.DataFrame = [value: string]

scala> val words = lines.as[String]
val words: org.apache.spark.sql.Dataset[String] = [value: string]

scala> val stopWords = Seq("is", "the", "a", "an", "and", "to")
val stopWords: Seq[String] = List(is, the, a, an, and, to)

scala> val cleaned = words.map { line =>
  val lower = line.toLowerCase
  val noSpace = lower.trim.replaceAll("\\s+", " ")
  val noPunct = noSpace.replaceAll("[a-zA-Z']+", "")
  val filtered = noPunct.split(" ").filter(word => !stopWords.contains(word))
  val lemmatized = filtered.map {
    case w if w.endsWith("ing") => w.dropRight(2)
    case w if w.endsWith("ed") => w.dropRight(2)
    case w => w
  }
  lemmatized.mkString(" ")
}
val cleaned: org.apache.spark.sql.Dataset[String] = [value: string]

scala> val query = cleaned.writeStream.outputMode("append").format("console").option("checkpointLocation", "C:/tmp/checkpoint").start()
val query: org.apache.spark.sql.streaming.StreamingQuery = org.apache.spark.sql.execution.streaming.runtime.StreamingQueryWrapper@53f6a2d4

-----
Batch: 0
-----
+-----+
|value|
+-----+

-----
Batch: 1
-----
+-----+
|          value|
+-----+
|the dance sing song|
+-----+

-----
Batch: 2
-----
+-----+
|          value|
+-----+
|cats are play in ...|
+-----+

-----
Batch: 3
-----
+-----+
|          value|
+-----+
|he walk talk teacher|
+-----+
```